



Texture synthesis, Style transfer

~852k parameters

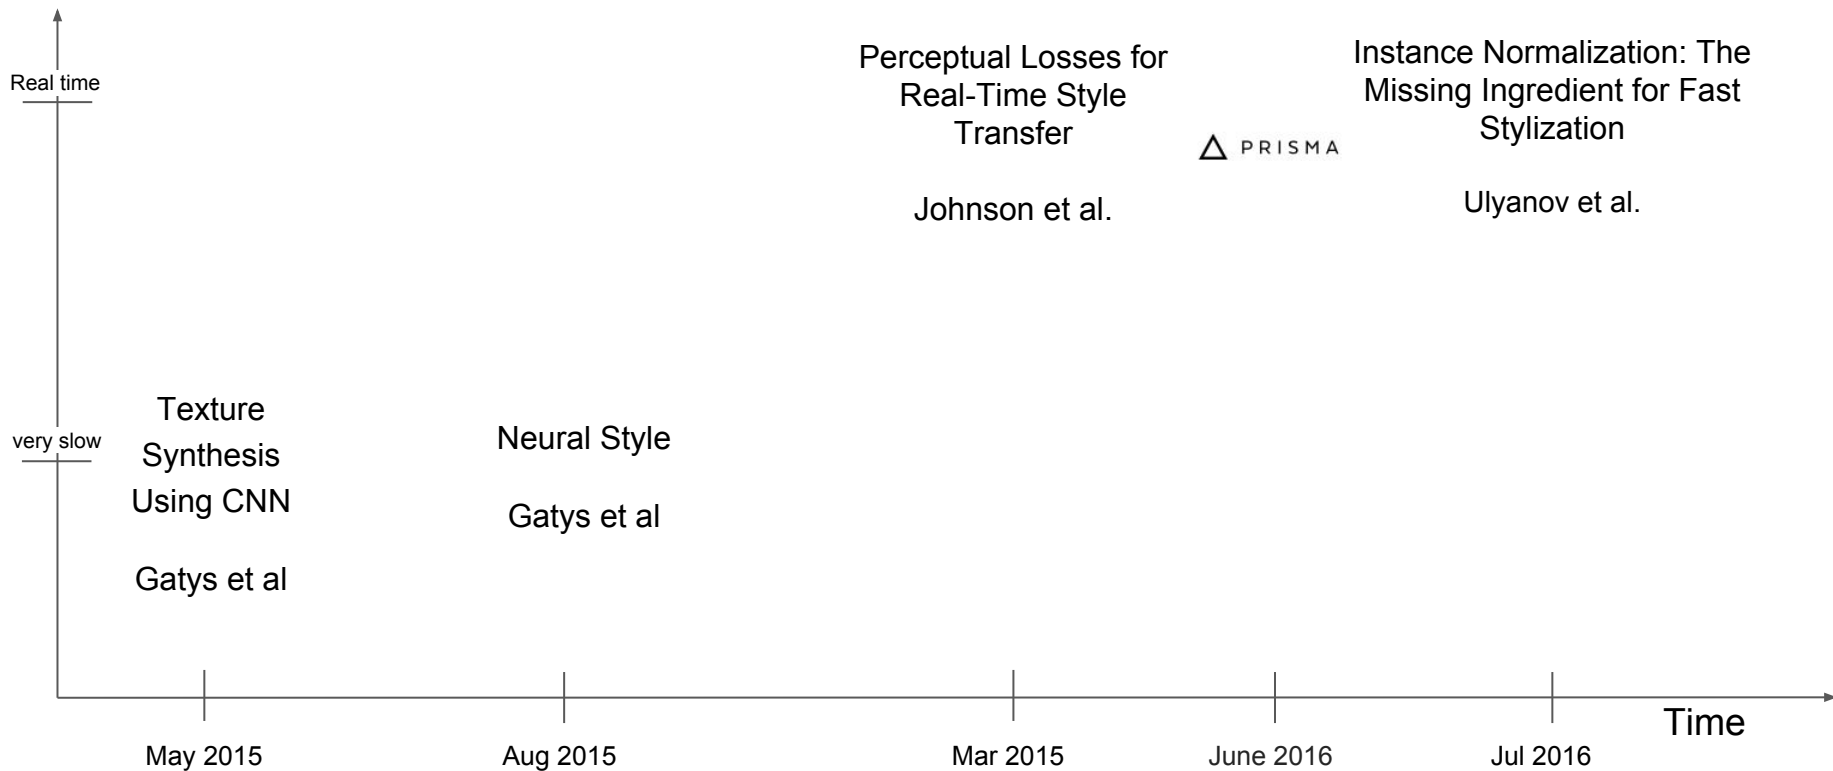


original





Big picture

Speed



Texture Synthesis Using CNN



Which one is
generated?

Texture Synthesis Using CNN


~852k parameters





original

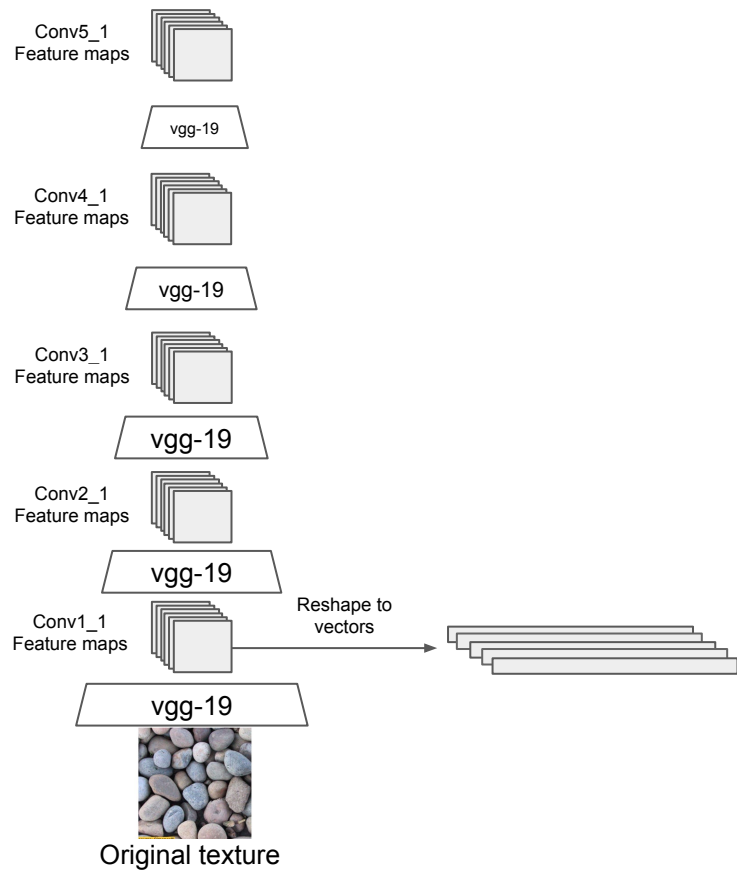


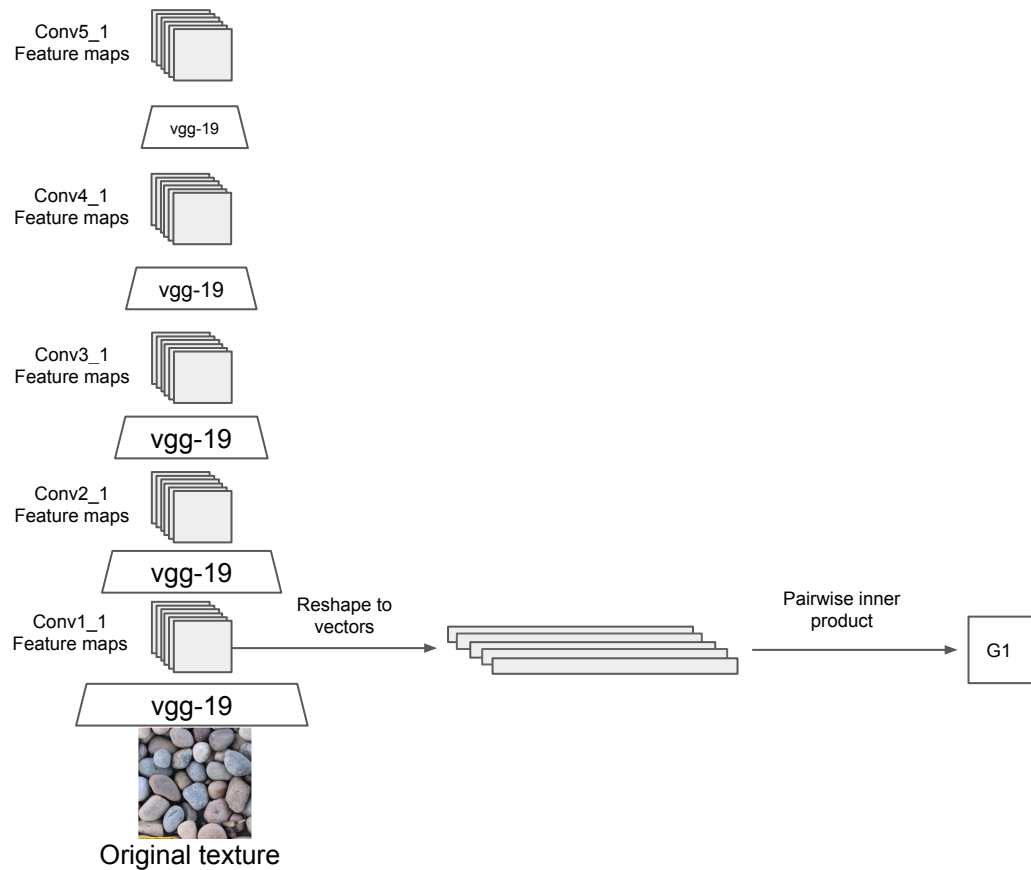
Which one is
generated?

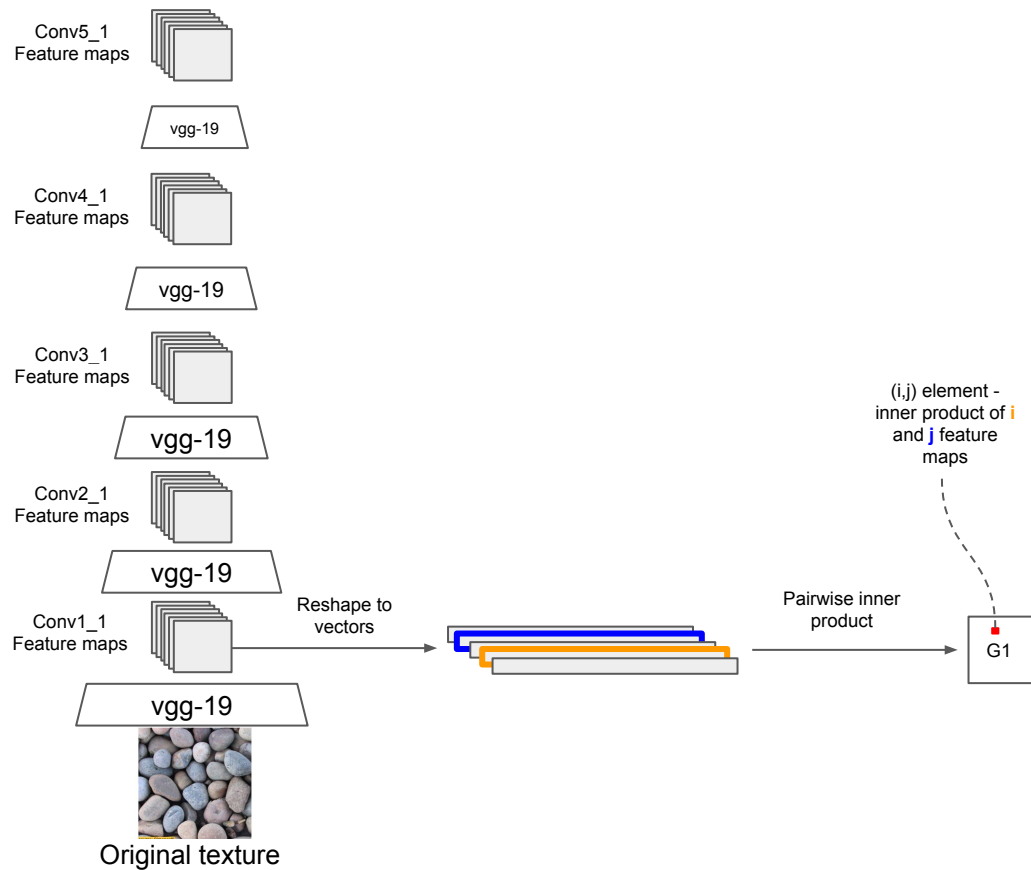


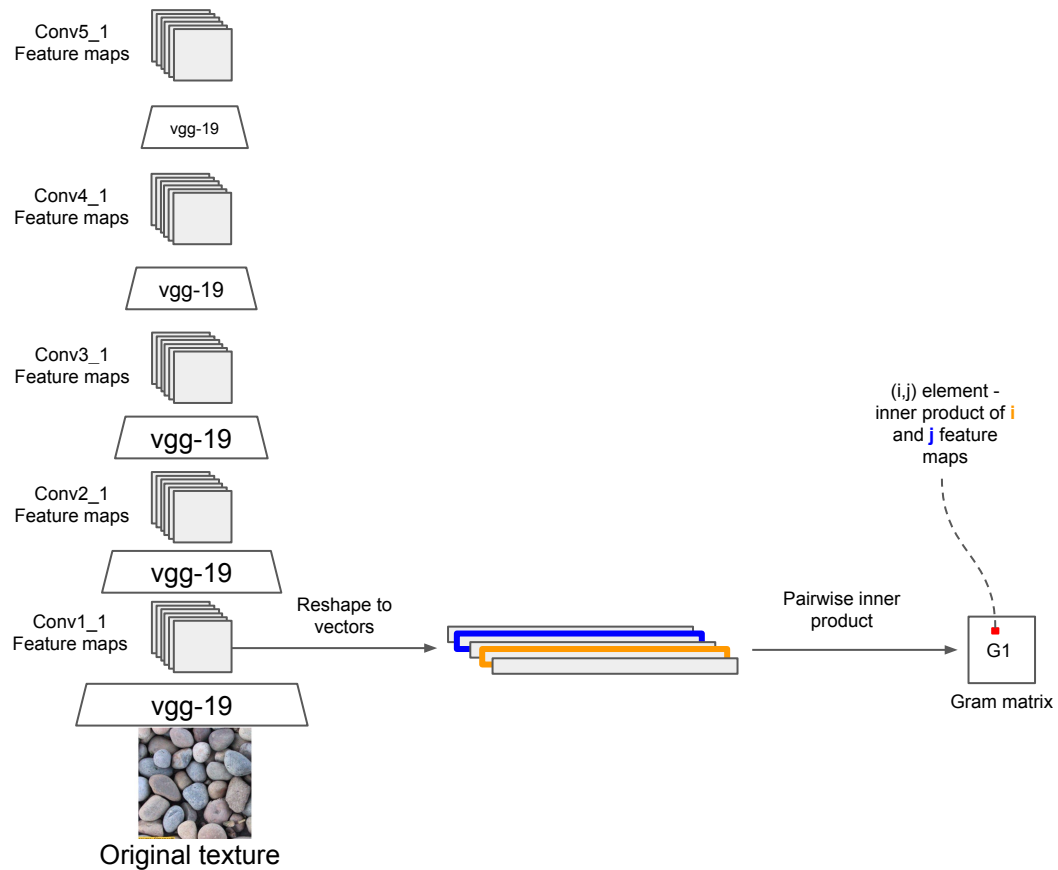





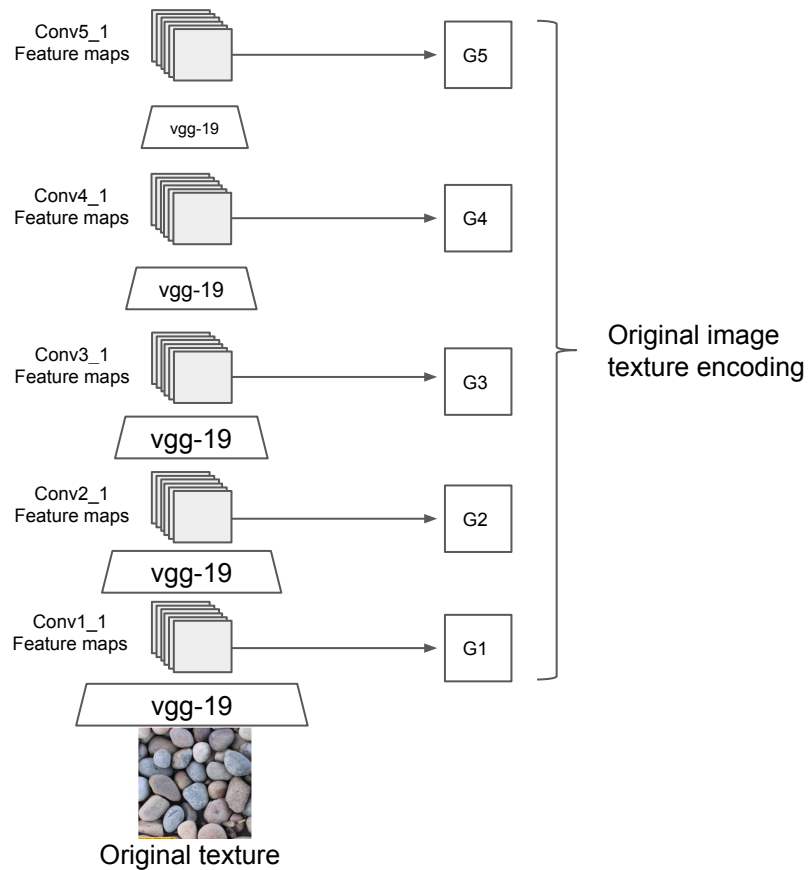




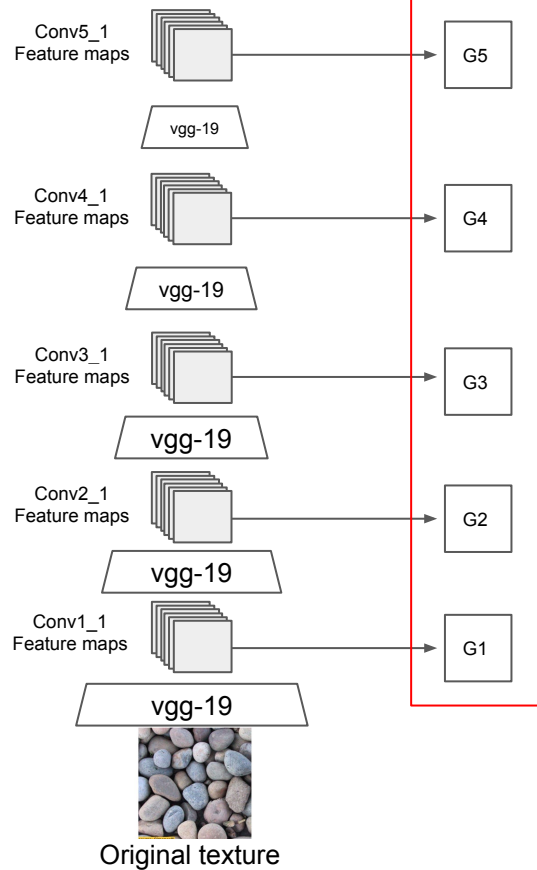






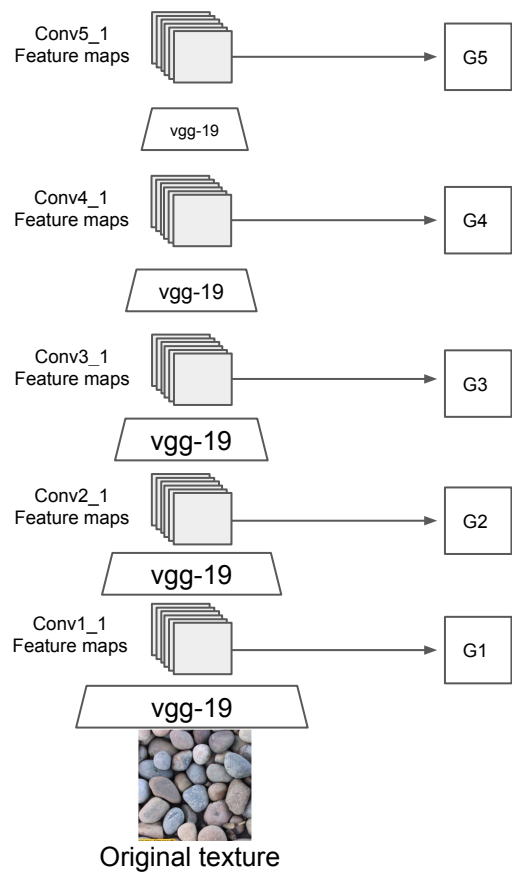


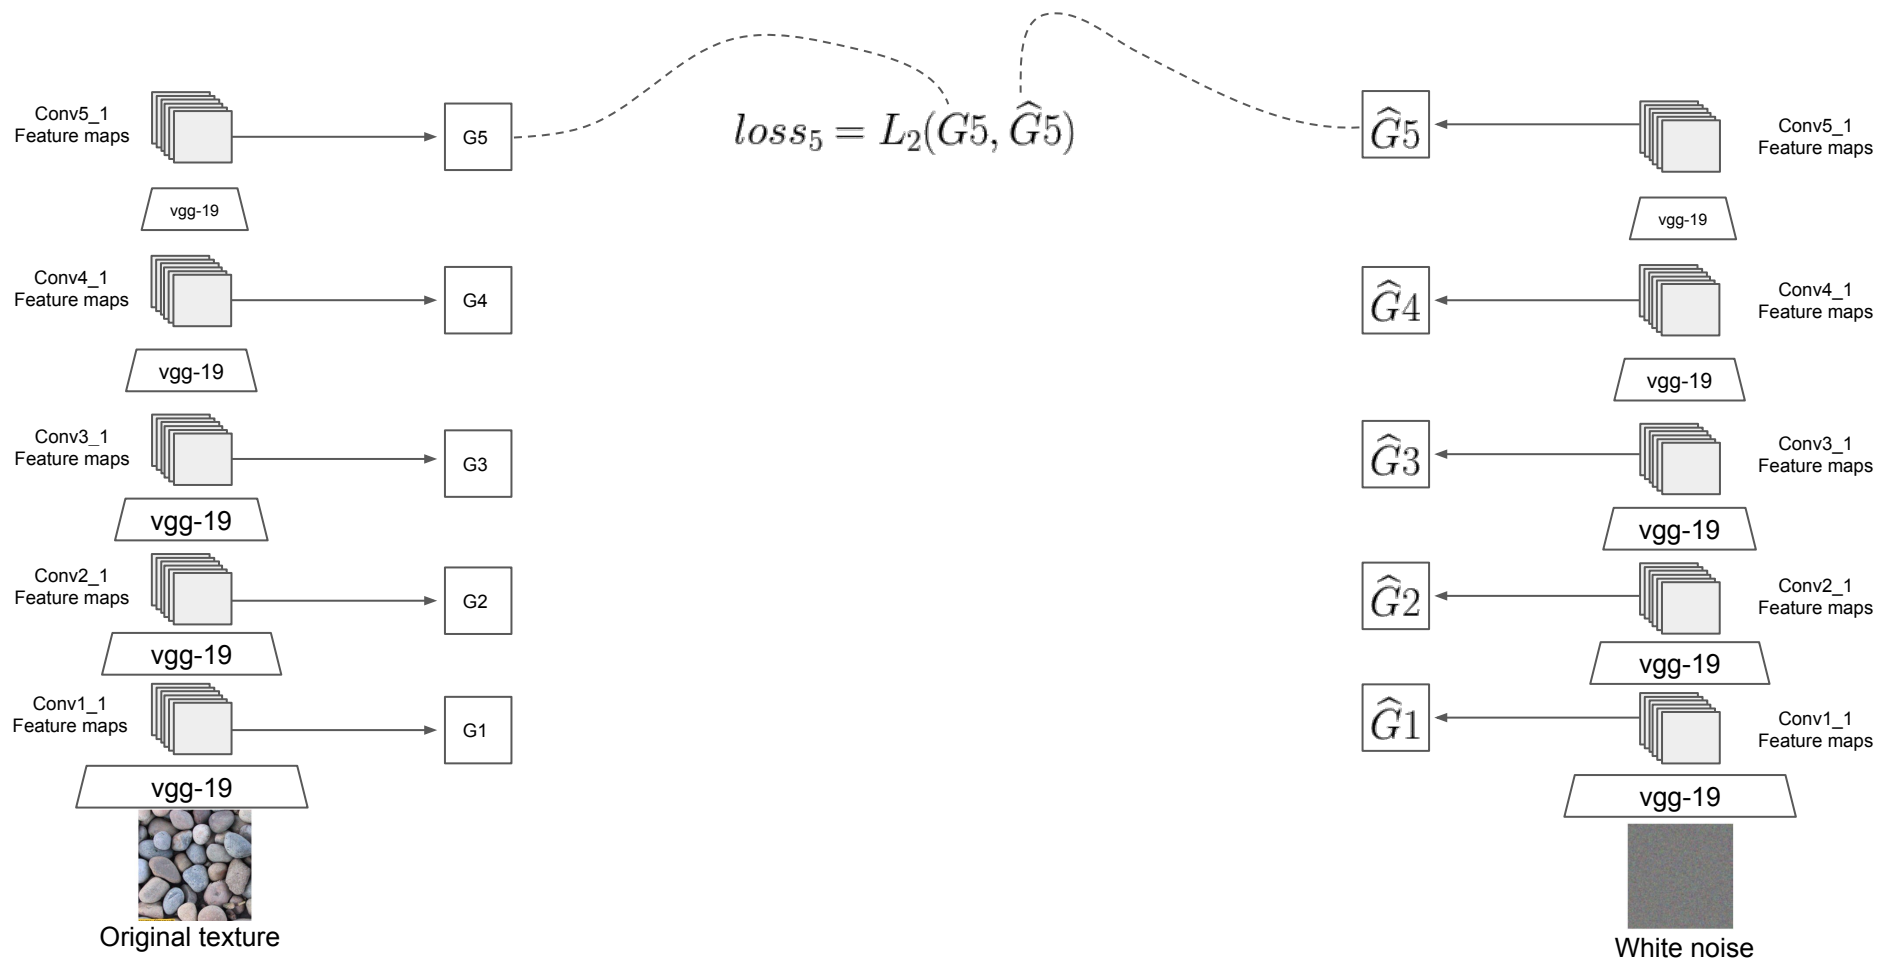


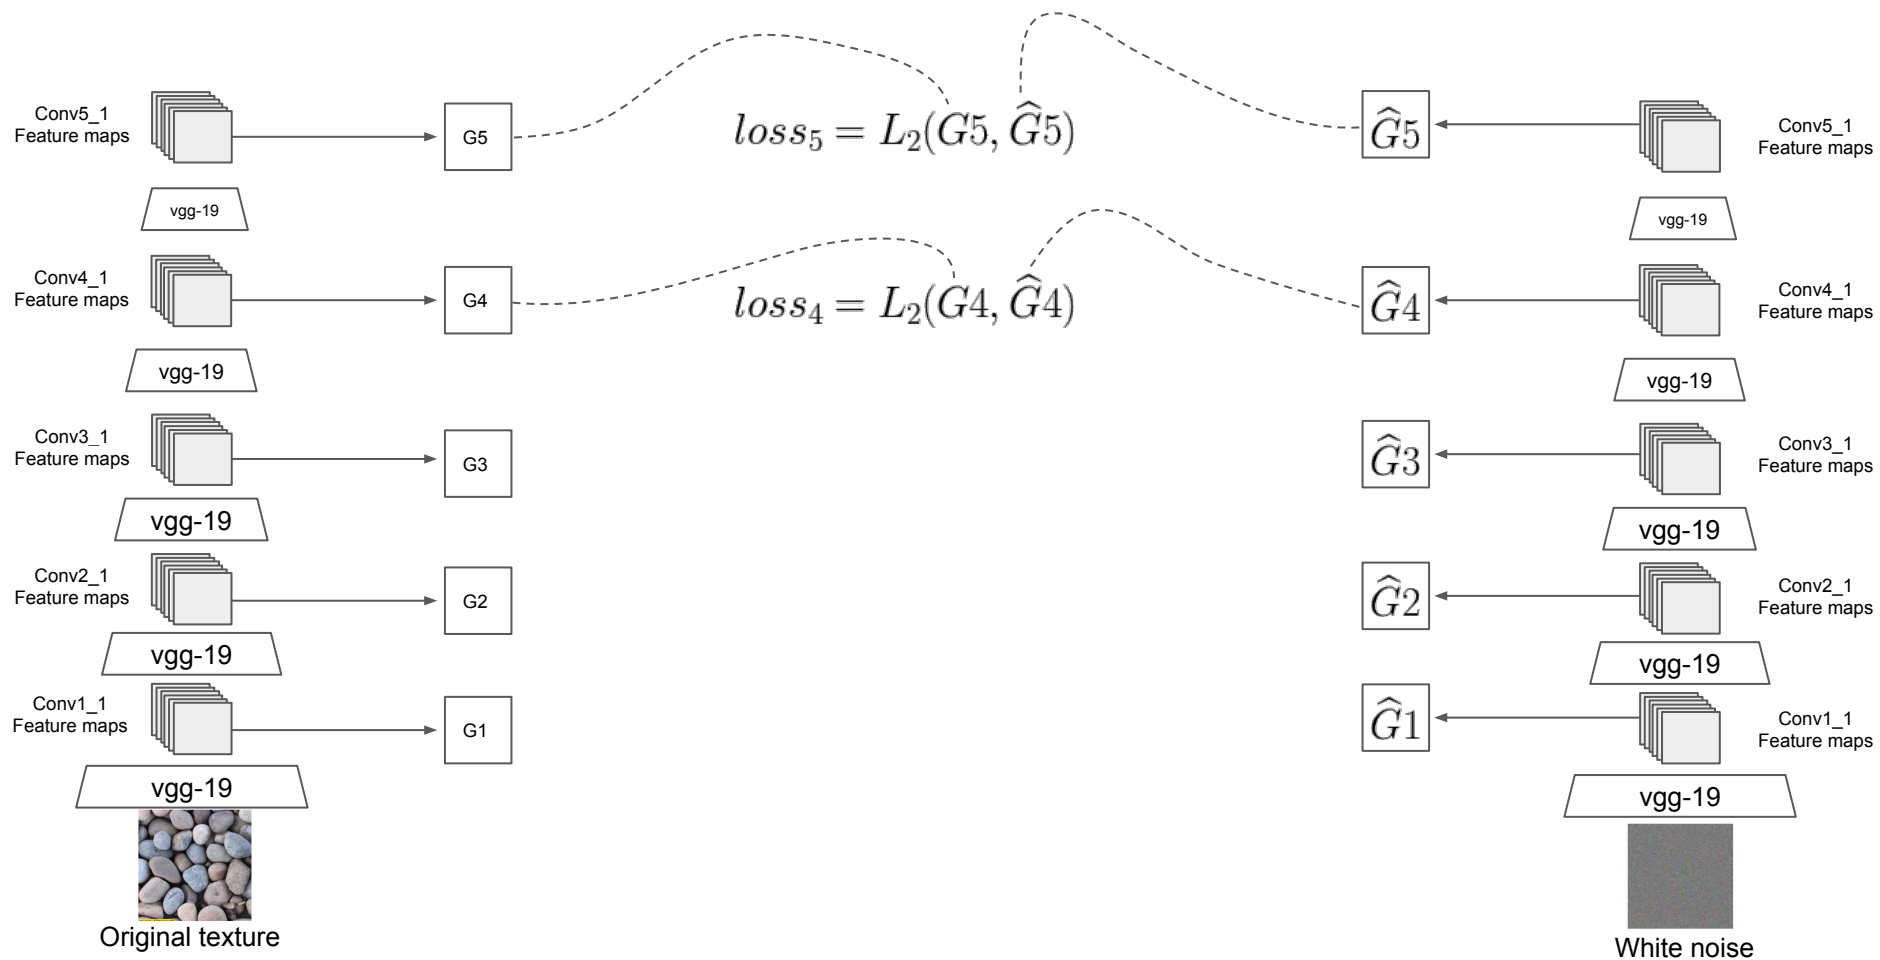
Save this

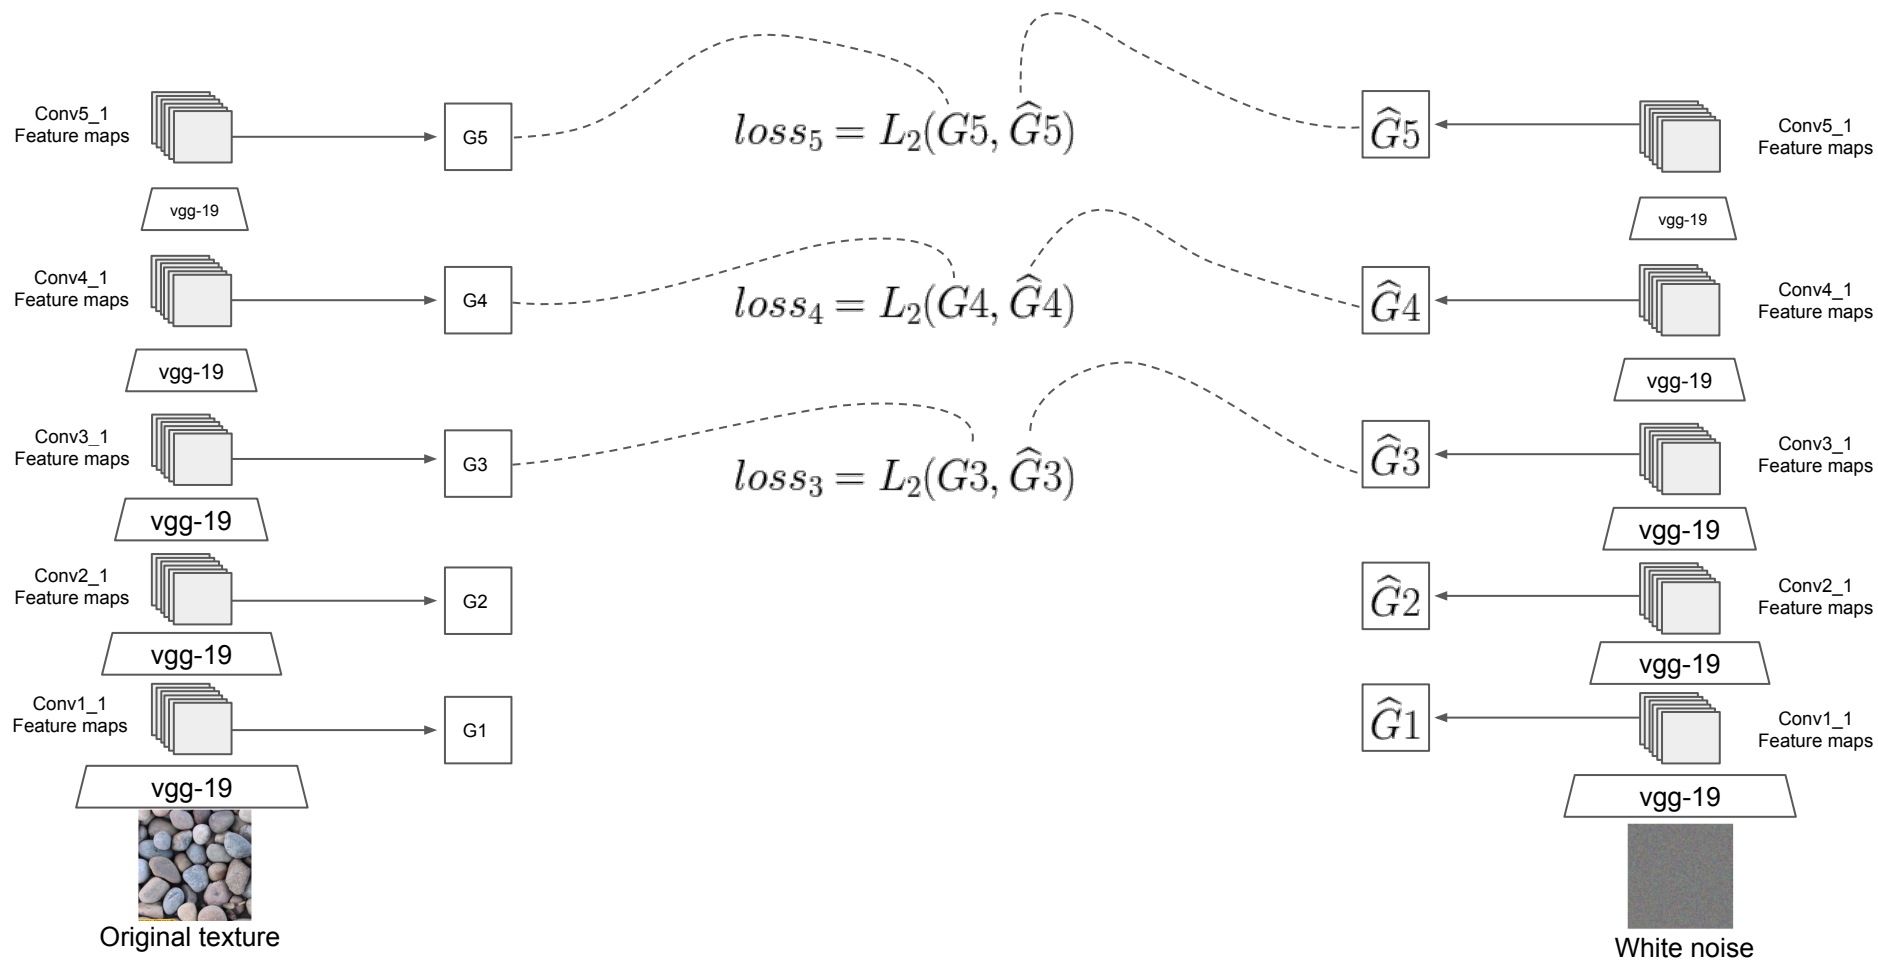


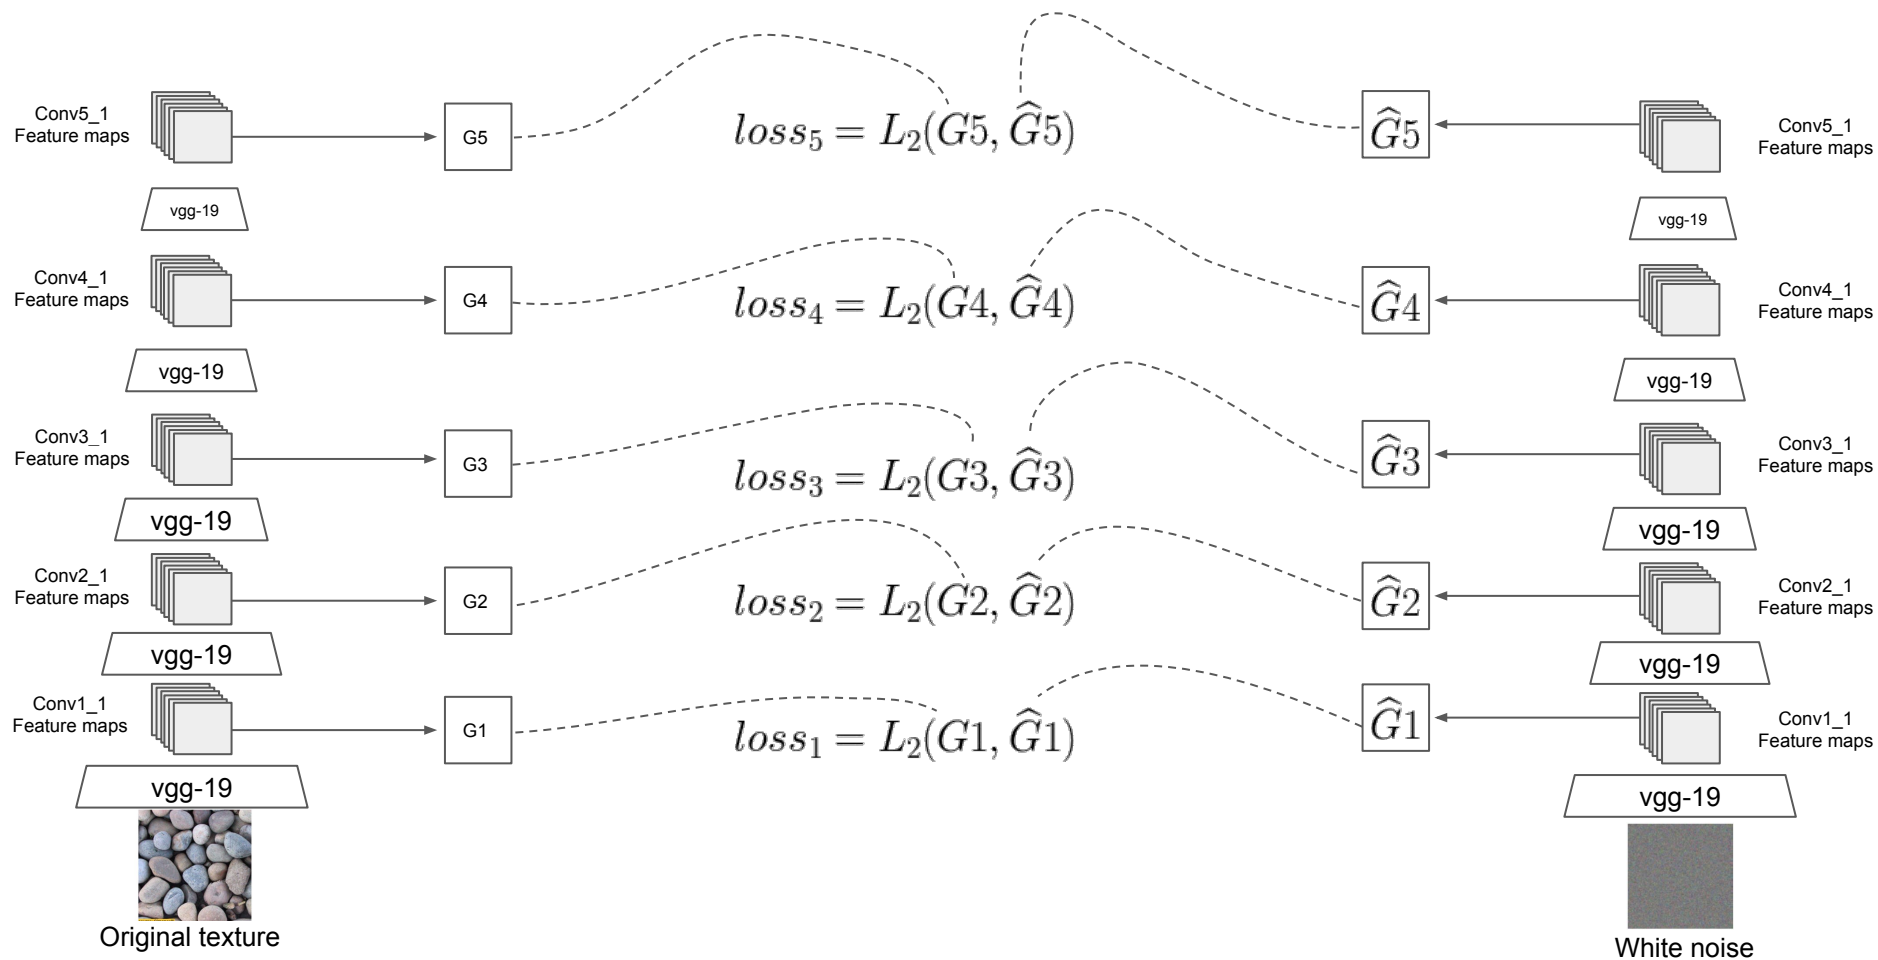


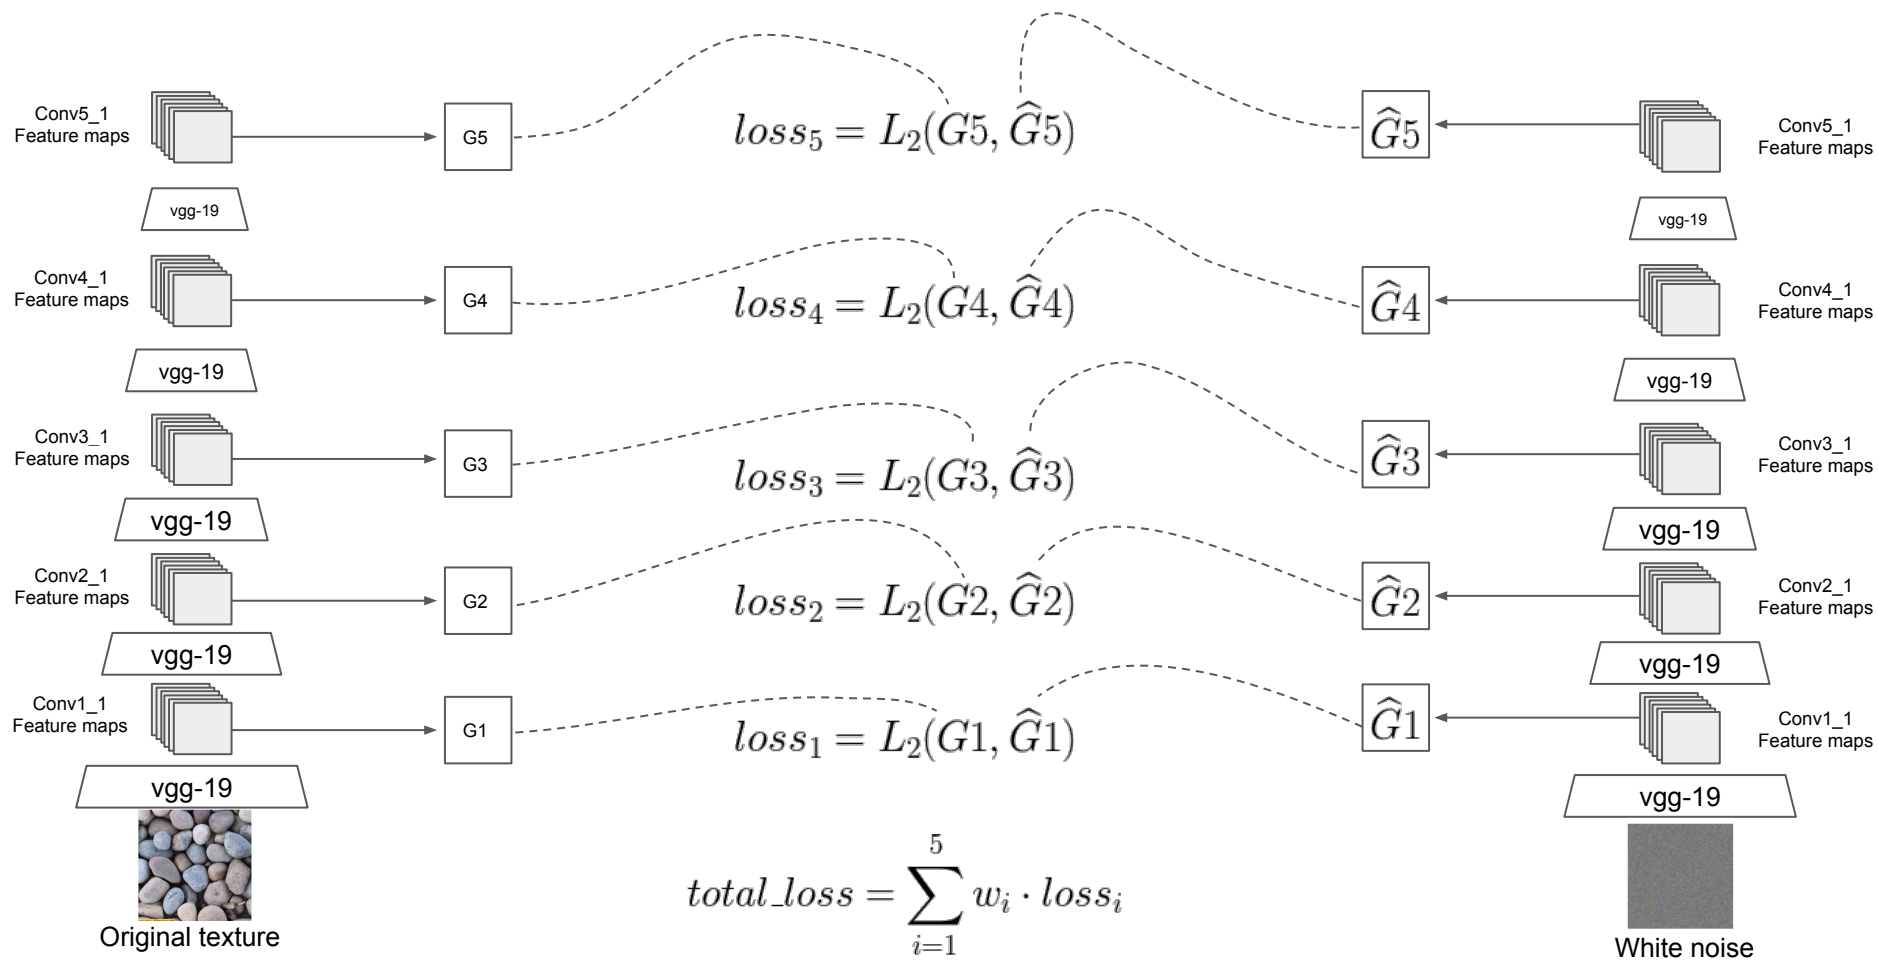












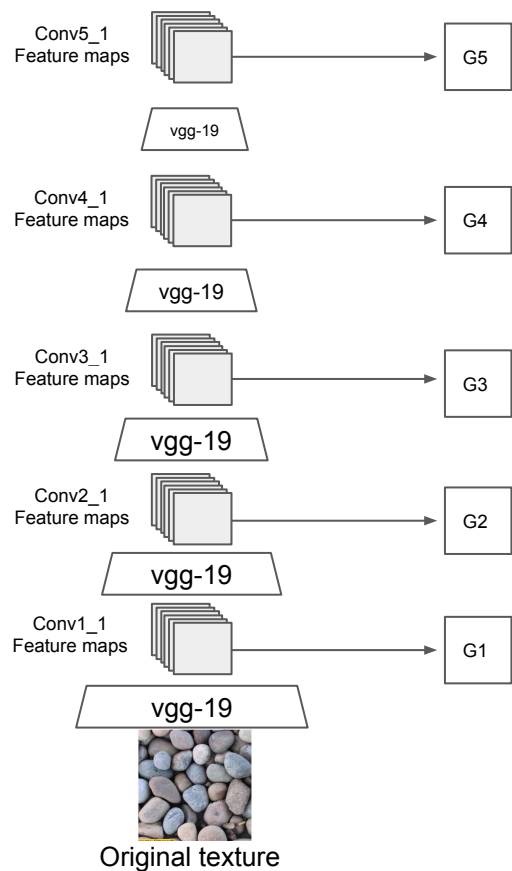




$$total_loss = \sum_{i=1}^5 w_i \cdot loss_i$$


Do backprop!











$$total_loss = \sum_{i=1}^5 w_i \cdot loss_i$$


Gradient for
image pixels












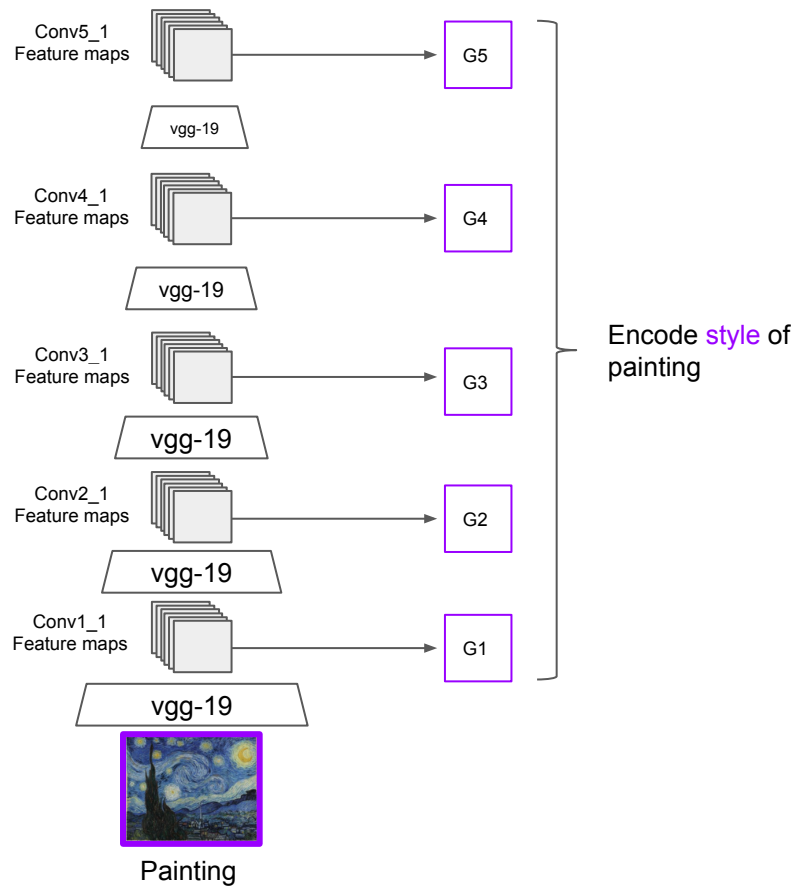
$$total_loss = \sum_{i=1}^5 w_i \cdot loss_i$$

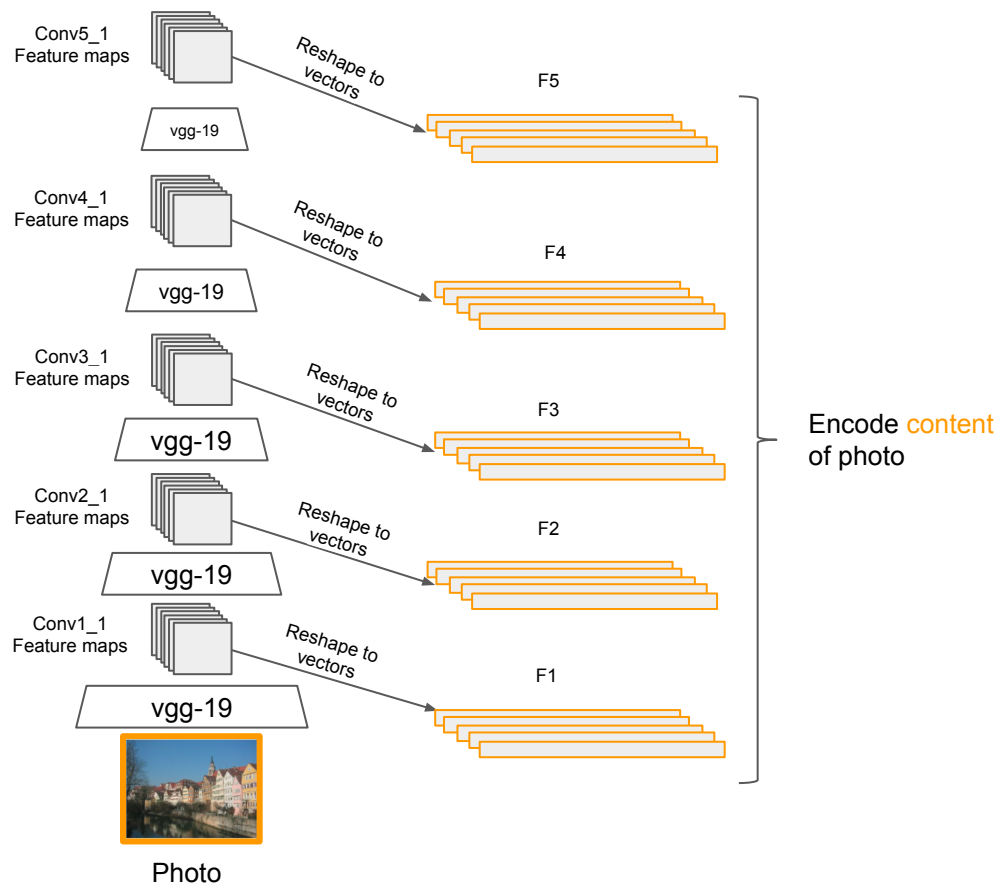
Repeat update
procedure few
times

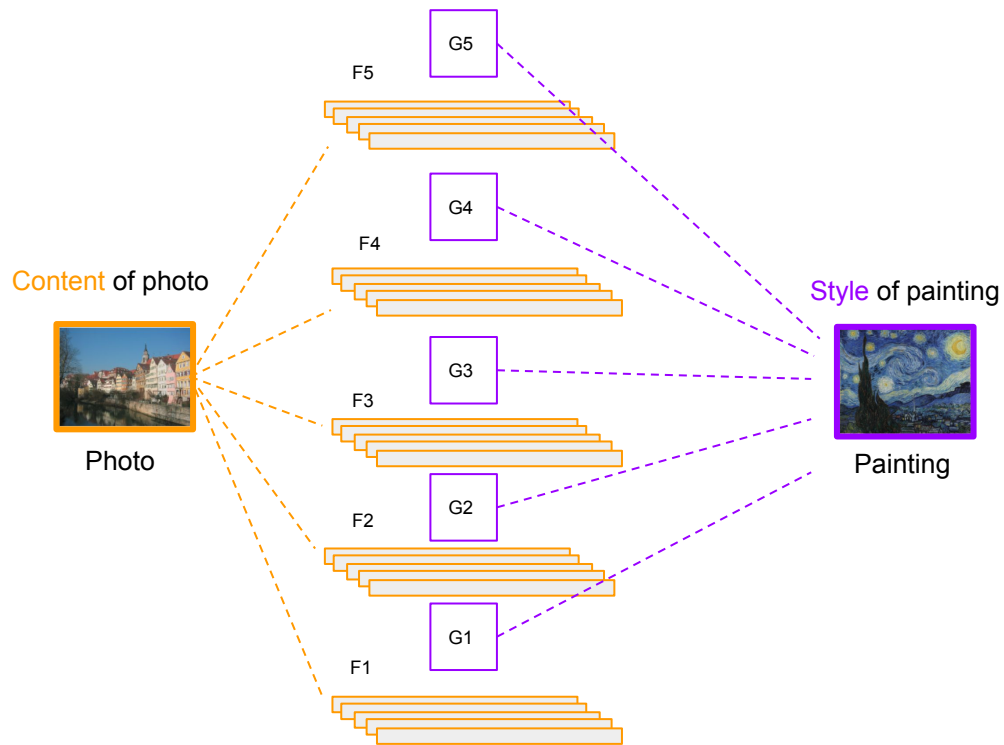





A neural algorithm for artistic style















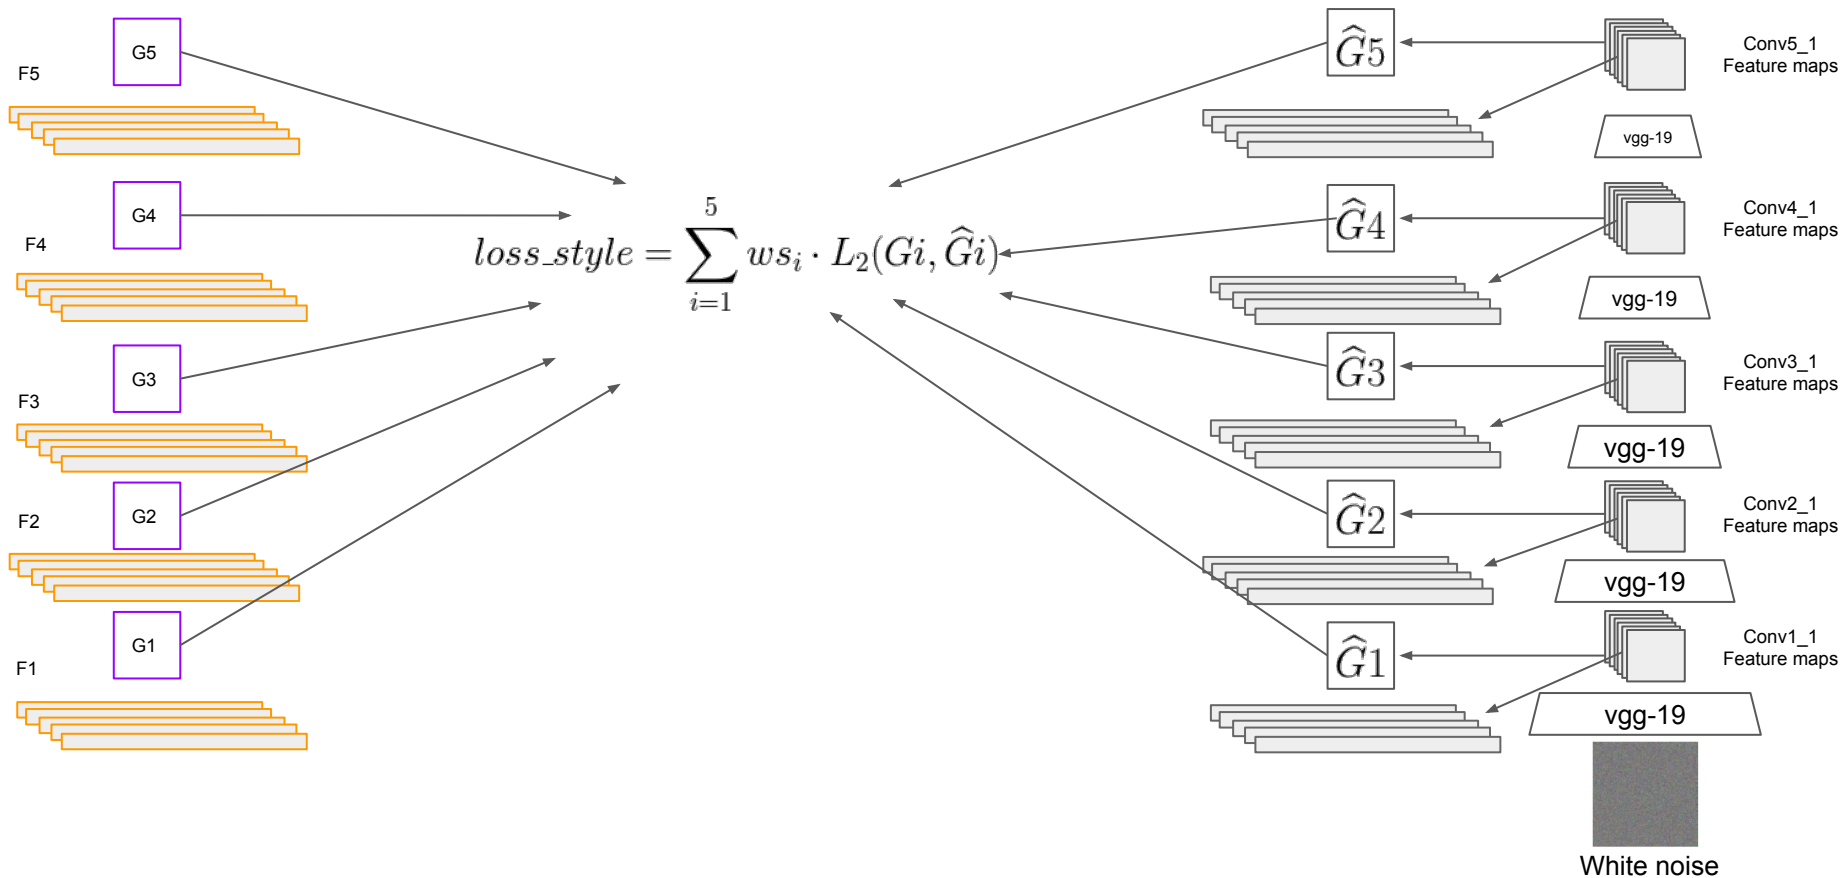


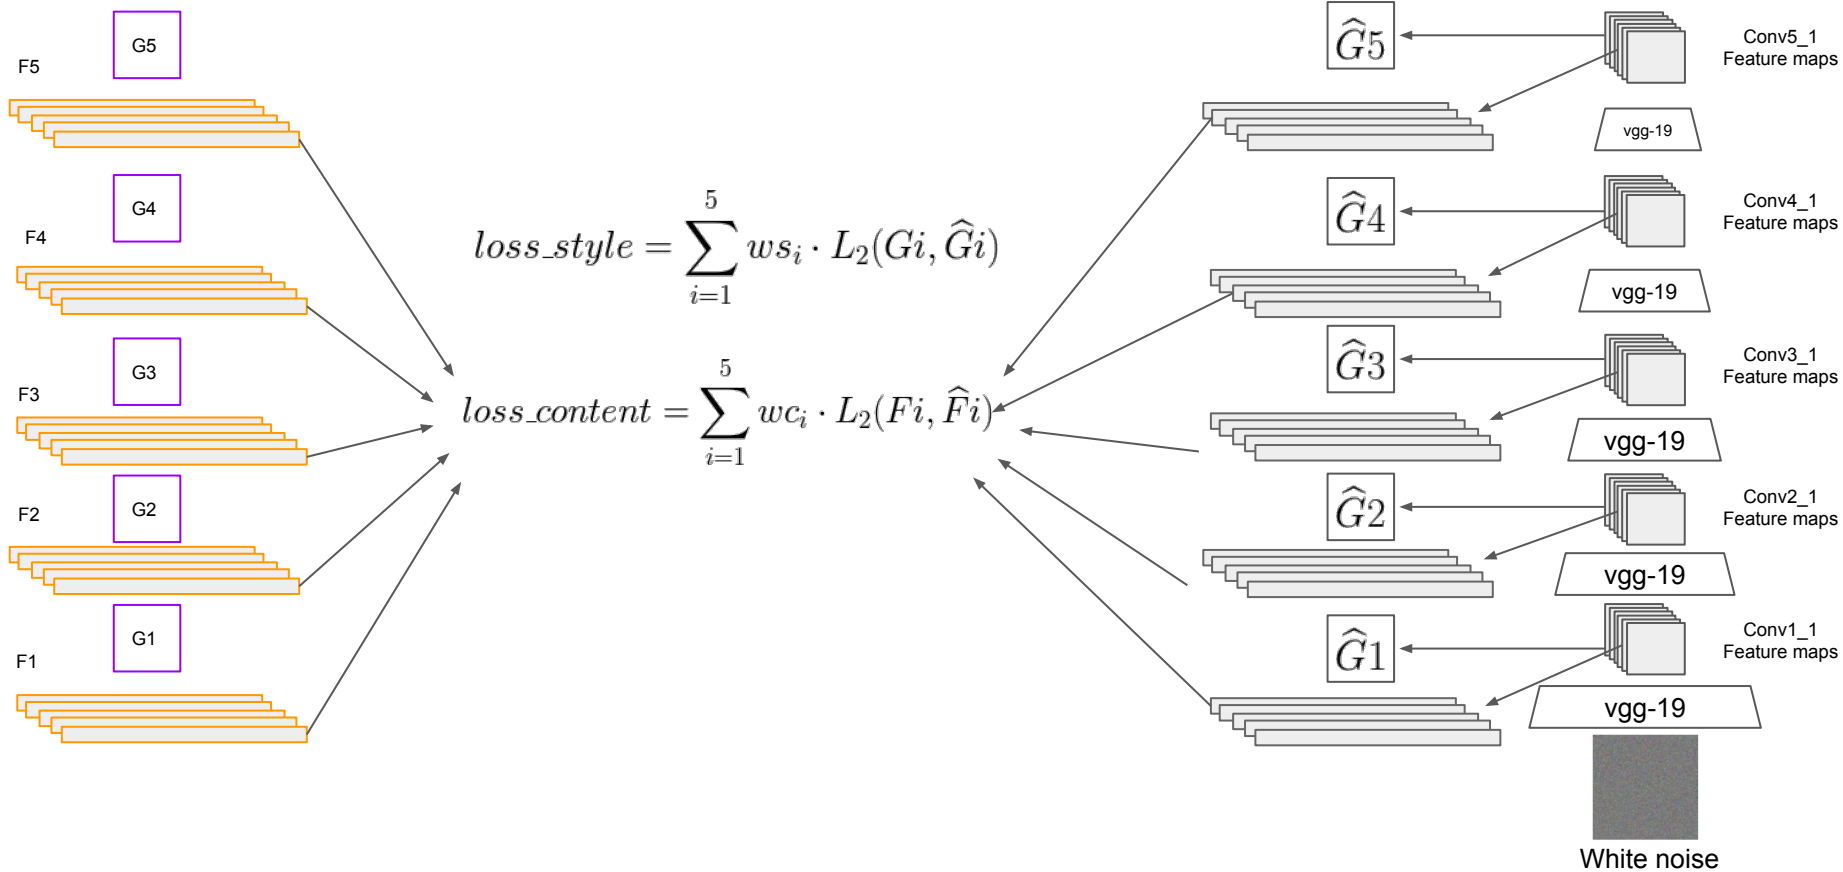
White noise
















$$loss_style = \sum_{i=1}^5 ws_i \cdot L_2(G_i, \hat{G}_i)$$

$$loss_content = \sum_{i=1}^5 wc_i \cdot L_2(F_i, \hat{F}_i)$$

$$total_loss = \alpha \cdot loss_content + \beta \cdot loss_style$$







$$loss_style = \sum_{i=1}^5 ws_i \cdot L_2(G_i, \hat{G}_i)$$

$$loss_content = \sum_{i=1}^5 wc_i \cdot L_2(F_i, \hat{F}_i)$$

$$total_loss = \alpha \cdot loss_content + \beta \cdot loss_style$$

Do backprop!






$$loss_style = \sum_{i=1}^5 ws_i \cdot L_2(G_i, \hat{G}_i)$$

$$loss_content = \sum_{i=1}^5 wc_i \cdot L_2(F_i, \hat{F}_i)$$


$$total_loss = \alpha \cdot loss_content + \beta \cdot loss_style$$



Gradient for
image pixels

$$\frac{\partial \mathcal{L}}{\partial \hat{x}}$$


White noise




$$loss_style = \sum_{i=1}^5 ws_i \cdot L_2(G_i, \hat{G}_i)$$

$$loss_content = \sum_{i=1}^5 wc_i \cdot L_2(F_i, \hat{F}_i)$$

$$total_loss = \alpha \cdot loss_content + \beta \cdot loss_style$$



Update only
image pixels




$$loss_style = \sum_{i=1}^5 ws_i \cdot L_2(G_i, \hat{G}_i)$$

$$loss_content = \sum_{i=1}^5 wc_i \cdot L_2(F_i, \hat{F}_i)$$



$$total_loss = \alpha \cdot loss_content + \beta \cdot loss_style$$

Repeat several
times




Intuition


Lets try to optimize the image (initially white noise) so that:



Intuition



Intuition





Key points

1. NeuralStyle was a start point for many research and commercial applications (prisma)
2. Non trivial example of transfer learning (deep learning is more than just classification or regression)
3. Object recognition model (VGG) is invariant to different object representations (style, shape, background, etc) => content and style can be divided

Links

- Arxiv (texture synthesis): <https://arxiv.org/abs/1505.07376>
- Arxiv (neural style): <https://arxiv.org/abs/1508.06576>
- Torch (neural style): <https://github.com/jcjohnson/neural-style>
- Torch (fast neural style): <https://github.com/jcjohnson/fast-neural-style>

Thank you!

Our website: deepsystems.io

Our team is looking for business partners to make exciting deep learning solutions.